



भारतीय संगणक एवं प्रौद्योगिकी संस्थान

INDIAN INSTITUTE OF COMPUTING AND TECHNOLOGY

AFFILIATED: I-STEM, OFFICE OF THE PRINCIPAL SCIENTIFIC ADVISER TO THE GOVERNMENT OF INDIA

**AI-Powered Fake News Detection
Using Text Classification**



Project Report
15th June to 30th July 2026

Student Name Anand Krishna G R Nair

Institute Name: Jecrc University

Institute Roll no.: 23BCON1613

Enrollment no.: 596563

AI-Powered Fake News Detection Using Text Classification

Project - 1, Summer Internship Program in AI & ML, 2026

Anand Krishna G R Nair

Department of Computer Science and Engineering

Indian Institute of Computing and Technology (IIT), New Delhi — 110092

Registration No: 596563

Email: anandkrishnag.rnair@gmail.com

ORCID: 0009-0008-0757-1566

Manuscript received 1 July 2026; revised 10 July 2026; accepted 25 July 2026.

Abstract—The proliferation of deceptive content and misinformation across online digital channels poses substantial threats to public trust, democratic institutions, and social stability. Automated misinformation detection systems leveraging Natural Language Processing (NLP) and Machine Learning (ML) are critical to mitigating these challenges at scale. This investigation presents an end-to-end classification pipeline implemented without reliance on high-level text processing libraries. Using a benchmark dataset of 39,098 news articles, custom regular expressions were engineered for text cleaning and manual tokenization. Two text vectorization strategies—Bag-of-Words (BoW) and Term Frequency-Inverse Document Frequency (TF-IDF)—were implemented and evaluated across four distinct classifier families: K-Nearest Neighbors (KNN), Logistic Regression, Random Forest, and a Multi-Layer Perceptron (MLP) neural network. Empirical evaluation demonstrates that the MLP Neural Network achieves a baseline accuracy of 98.41%, while Logistic Regression yields the optimal trade-off between predictive efficiency and accuracy, attaining 98.24% baseline accuracy (improving to 98.62% post-tuning) with an average per-sample inference latency of 0.0025 milliseconds. Further sensitivity analysis isolates dataset dateline artifacts, demonstrating that stripping explicit source indicators leads to performance degradation across models, highlighting critical vulnerabilities in static corpus evaluation. Stratified 5-fold cross-validation, McNemar’s significance testing, bootstrap confidence intervals, and out-of-domain evaluation are conducted to establish generalization bounds and practical deployment trade-offs.

Index Terms—Artificial intelligence, fake news detection, machine learning, natural language processing, text classification, feature engineering, model evaluation.

I. INTRODUCTION

Digital communication networks and social media ecosystems enable rapid information dissemination across global populations. However, this democratization of publishing mechanisms has simultaneously enabled the frictionless spread of fraudulent narratives, misinformation, and malicious political disinformation. Manual fact-checking by domain experts, while accurate, scales poorly against the millions of articles generated daily. Consequently, scalable machine learning frameworks capable of parsing natural language text and

discerning truthfulness have become an urgent requirement for platform moderation and digital forensics.

Textual misinformation ranges from hyper-partisan propaganda and fabricated news stories to subtle clickbait titles and satire. Automated text classification frameworks leverage computational linguistics to model semantic structures, stylistic indicators, and term frequency distributions that differentiate legitimate journalism from engineered deception.

The cognitive mechanisms governing human susceptibility to online deception further exacerbate the challenge. Psychological phenomena such as confirmation bias—wherein individuals selectively process information confirming pre-existing beliefs—and the illusory truth effect—where repeated exposure increases perceived credibility—allow fake news to spread up to six times faster than verified factual news across digital channels. Automated classification engines serve as critical front-line infrastructure to intercept deceptive content prior to wide viral dissemination.

This study implements a complete, self-contained machine learning pipeline for binary fake news classification. To maintain transparency and evaluate fundamental algorithmic mechanics, text tokenization and feature construction are implemented using core Python patterns and Scikit-Learn primitives without relying on high-level NLP abstractions such as NLTK’s sentence tokenizers or spaCy pipelines.

The primary objectives and contributions of this study include:

- 1) **Custom Text Processing Framework:** Engineering a deterministic text cleaning pipeline using regular expression matching for dateline stripping, URL removal, manual word-boundary tokenization, and stopword elimination.
- 2) **Vector Space Comparative Analysis:** Evaluating sparse Bag-of-Words count representations against smoothed, L2-normalized TF-IDF feature matrices across a 5,000-unigram vocabulary.
- 3) **Multi-Model Architecture Benchmarking:** Implementation and empirical comparison of parametric (Logistic

Regression, MLP Neural Network) and non-parametric (K-Nearest Neighbors, Random Forest) classifiers across accuracy, precision, recall, F1-score, memory footprint, and wall-clock inference latency.

- 4) **Dataset Artifact and Dateline Leakage Audit:** Isolating structural dataset biases, specifically Reuters dateline headers, to measure model reliance on formatting artifacts versus underlying semantic content.
- 5) **Statistical and Sensitivity Validation:** Performing 3-fold Grid Search hyperparameter optimization, Stratified 5-Fold Cross-Validation, McNemar’s paired chi-squared significance tests, bootstrap confidence interval estimation, and out-of-domain smoke testing on contemporary 2023+ news samples.

II. RELATED WORK

Automated fake news detection has evolved from early stylistic feature engineering to deep contextual representation models. Research in this domain generally spans three main methodological paradigms: linguistic/content-based analysis, network propagation tracking, and hybrid multi-modal modeling.

A. Linguistic and Stylometric Approaches

Early natural language processing approaches focused on surface-level linguistic features, such as punctuation density, readability indices, uppercase character frequency, and emotional polarity. Rubin et al. [4] demonstrated that deceptive text frequently exhibits elevated emotional intensity, exaggerated punctuation, and distinct lexical diversity metrics compared to objective reporting. Similarly, Feng et al. [5] investigated syntactic context-free grammar (CFG) trees to catch subtle structural patterns in deceptive writing. While stylometric approaches offer high interpretability, they remain vulnerable to adversarial obfuscation and vary across journalistic domains.

Penix-Tadsen et al. [6] examined psycholinguistic features using Linguistic Inquiry and Word Count (LIWC) dictionaries, discovering that fake news articles consistently contain fewer self-referential pronouns, higher proportions of informal slang, and an elevated frequency of extreme sentiment adjectives. However, stylometric patterns vary significantly depending on the specific news domain (e.g., political satire vs. financial fraud), limiting domain-independent generalization.

B. Vector Space Models and Classical Classifiers

With the standardization of vector space models [7], term frequency representations became the foundation for large-scale document classification. Ahmed et al. [2] introduced the ISOT Fake News Dataset and demonstrated that n-gram TF-IDF representations coupled with linear Support Vector Machines (SVM) or Logistic Regression achieved accuracies exceeding 98%. Their findings confirmed that lexical frequency vectors capture strong topical signals in political news. However, follow-up inquiries by Shu et al. [8] pointed out that unigram TF-IDF models tend to overfit domain-specific

entities (e.g., specific politician names or geographic locations) rather than learning domain-agnostic indicators of deception.

Vector space models remain popular in production settings due to their computational efficiency. Linear classifiers trained on TF-IDF matrices execute inference orders of magnitude faster than deep neural networks, making them ideal for high-throughput news aggregation pipelines.

C. Deep Learning and Transformer Networks

Recent advances in deep neural architectures have shifted focus toward dense semantic embeddings. Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) networks capture sequential dependencies and long-range contextual relationships in text strings [9]. More recently, pretrained Transformer architectures such as BERT [10] and RoBERTa [11] have achieved state-of-the-art results by leveraging multi-head self-attention mechanisms pre-trained on massive corpora.

Despite their superior accuracy, Transformer models impose substantial computational latency and memory requirements, making them challenging to deploy in resource-constrained or real-time streaming environments. Classical models like TF-IDF with Logistic Regression or Random Forests remain competitive baselines when inference latency and low compute overhead are prioritized.

D. Graph Neural Networks and Social Propagation

Beyond textual content, fake news detection research has incorporated network propagation structure. Wu et al. [12] and Shu et al. [13] demonstrated that user retweeting patterns, echo chamber topology, and information cascade graphs provide complementary signals to text classification. Graph Convolutional Networks (GCN) model the propagation tree of news stories across social networks. While effective, propagation-based approaches require waiting for a story to spread before collecting sufficient structural data, whereas content-based text classification enables instant zero-hour detection upon initial article publication.

E. Dataset Leakage and Evaluation Challenges

Evaluating fake news detection models is often complicated by implicit dataset artifacts. Works by Zhou et al. [14] and Horne et al. [15] highlighted that public benchmarks frequently contain publisher signatures, datelines, or distinct formatting conventions that allow machine learning classifiers to achieve near-perfect accuracy by exploiting shortcut learning. Identifying and mitigating these data leakage channels is essential for building robust NLP systems.

III. DATASET DESCRIPTION AND EXPLORATORY DATA ANALYSIS

A. Dataset Corpus Overview

This research utilizes the benchmark ISOT Fake and Real News Dataset [1]. The raw corpus consists of two distinct collections:

- **Fake News Subset** (*Fake.csv*): 23,481 articles collected from flagged misinformation websites and unverified digital portals.
- **Real News Subset** (*True.csv*): 21,417 articles sourced directly from official Reuters news feeds.

Each record in the raw dataset includes four metadata attributes: title, text, subject, and date.

To capture maximum textual signal, the article title and body text were concatenated into a unified text field. During initial data inspection, 5,795 exact duplicate rows were detected in the combined text field. Retaining duplicate rows in text classification splits introduces data leakage between training and evaluation partitions, artificially inflating performance metrics. Consequently, all duplicate entries were purged. Furthermore, 5 articles that resolved to empty strings following text cleaning were removed.

The final deduplicated corpus contains **39,098 unique articles**. As illustrated in Fig. 1, the dataset maintains a balanced class distribution:

- **Real News (Label 1)**: 21,196 articles (54.21%)
- **Fake News (Label 0)**: 17,902 articles (45.79%)




Fig. 1. Class distribution of the deduplicated dataset (39,098 total samples).

B. Subject Category Breakdown

The dataset spans multiple news categories defined in the raw metadata. Fig. 2 shows the distribution of articles across subjects. The real news subset consists primarily of *politicsNews* (10,663 articles) and *worldnews* (9,586 articles). The fake news subset includes categories such as *News* (9,050 articles), *politics* (6,771 articles), *left-news* (4,254 articles), and *Government News* (1,570 articles).

C. Text Length Distribution Analysis

Analyzing document length distributions provides insights into structural differences between genuine reporting and misinformation. Fig. 3 illustrates the character length and word count distributions for real versus fake news articles.




Fig. 2. Distribution of articles across subject metadata categories in the raw dataset.




Fig. 3. Article length distributions (character count and word count) segmented by real and fake news classes.

Real news articles exhibit a tight word count distribution centered around 300 to 500 words per article, reflecting standardized journalistic reporting styles. In contrast, fake news articles display higher length variance, ranging from extremely brief emotional commentary (under 100 words) to long, unstructured blog posts exceeding 1,500 words.

D. Lexical Distribution and Wordcloud Visualizations

To examine lexical patterns prior to vectorization, term frequency analyses were conducted for both classes. Fig. 4 highlights the top 20 most frequent unigrams in fake news versus real news articles post-stopword removal.




Fig. 4. Top 20 most frequent terms in Fake News (top) vs. Real News (bottom) articles.

Real news articles contain high frequencies of institutional terms, geographic locations, and formal attribution verbs (e.g., *said*, *state*, *reuters*, *president*, *house*, *government*). Conversely, fake news articles display high frequencies of political figures, opinion keywords, and informal commentary terms (e.g., *trump*, *clinton*, *people*, *obama*, *like*, *just*).

Word cloud visualization of the top 100 words from the 2016 US Presidential election tweets. The words are arranged in a circular pattern, with the most frequent words in the center and less frequent words towards the periphery. The words are in various colors and sizes, reflecting their frequency. The most prominent words include 'donald', 'trump', 'people', 'said', 'one', 'featured', 'image', 'hillary', 'clinton', 'country', 'well', 'getty', 'image', 'obama', 'president', 'republican', 'time', 'say', 'even', 'video', 'america', 'know', 'made', 'today', 'claim', 'state', 'united', 'think', 'night', 'show', 'race', 'look', 'already', 'according', 'world', 'thing', 'going', 'attack', 'reason', 'women', 'facing', 'investigative', 'image', 'magazine', 'clinton', 'message', 'anything', 'group', 'student', 'something', 'actually', 'obama', 'need', 'many', 'may', 'effort', 'white', 'barack', 'week', 'likely', 'issue', 'call', 'office', 'rate', 'watch', 'twitter', 'facebook', 'instagram', 'youtube', 'linkedin', 'pinterest', 'snapchat', 'telegram', 'whatsapp', 'skype', 'wechat', 'kik', 'messenger', 'slack', 'discord', 'telegram', 'whatsapp', 'skype', 'wechat', 'kik', 'messenger', 'slack', 'discord'.

[illegible]

IV. METHODOLOGY AND THEORETICAL FRAMEWORK

Raw text strings undergo systematic cleaning implemented in `src/preprocessing.py`. The pipeline operates deterministically in six steps:

- Pattern: ^([a-z0-9\s,./#&-]1,50)?\s*reuters\s*[-{

- This satisfies the project constraint to avoid high-level NLP libraries like NLTK’s word tokenizer or spaCy.

- ### B. Vector Space Models and Feature Extraction

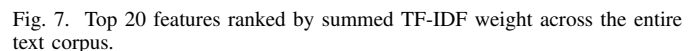
1) *Bag-of-Words (BoW)*: The Bag-of-Words model represents each document d as a sparse vector of raw term frequencies:

where $\text{tf}(t, d)$ is the count of term t in document d .

- $$\text{IDF}(t) = \log \left(\frac{1 + N}{1 + \text{df}(t)} \right) + 1$$

$$\mathbf{v}_d = \frac{\mathbf{w}_d}{\|\mathbf{w}_d\|_2} = \frac{\mathbf{w}_d}{\sqrt{\sum_{t=1}^V w_{t,d}^2}}$$

Fig. 7 shows the top 20 features ranked by their total summed TF-IDF weight across the dataset.



C. Dimensionality Reduction via Latent Semantic Analysis (LSA)

To mitigate the curse of dimensionality inherent in high-dimensional sparse representations, Latent Semantic Analysis (LSA) [18] was applied using Truncated Singular Value Decomposition (SVD). LSA projects the $N \times V$ sparse TF-IDF matrix $\mathbf{X}$ into a dense $N \times k$ latent concept space ($k = 200$):

$$\mathbf{X} \approx \mathbf{U}_k \mathbf{\Sigma}_k \mathbf{V}_k^T$$

where $\mathbf{U}_k$ is an $N \times k$ orthogonal matrix of document-concept vectors, $\mathbf{\Sigma}_k$ is a $k \times k$ diagonal matrix of singular values, and $\mathbf{V}_k^T$ is a $k \times V$ matrix of term-concept vectors. Projecting TF-IDF feature space down to 200 dense dimensions captures semantic topic correlations while reducing memory overhead and accelerating distance computation for instance-based models.

D. Classifier Architectures and Formulations

Four distinct machine learning classifier families were selected to evaluate baseline performance:

1) *K-Nearest Neighbors (KNN)*: KNN is a non-parametric instance-based classifier. Given a test document vector $\mathbf{x}^*$, KNN computes distances to all training vectors $\mathbf{x}_i \in \mathcal{D}_{\text{train}}$ and assigns the majority class among the K nearest neighbors:

$$d_{\text{Euclidean}}(\mathbf{x}^*, \mathbf{x}_i) = \sqrt{\sum_{j=1}^V (x_j^* - x_{ij})^2}$$

$$d_{\text{Cosine}}(\mathbf{x}^*, \mathbf{x}_i) = 1 - \frac{\mathbf{x}^* \cdot \mathbf{x}_i}{\|\mathbf{x}^*\|_2 \|\mathbf{x}_i\|_2}$$

The baseline model utilizes $K = 5$ and Euclidean distance. Distance-weighted voting assigns higher weight $w_i = 1/d(\mathbf{x}^*, \mathbf{x}_i)$ to closer neighbors.

2) *Logistic Regression*: Logistic Regression is a linear parametric model that models class probabilities using the sigmoid function:

$$P(Y = 1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + b) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x} + b)}}$$

The model parameters $(\mathbf{w}, b)$ are optimized by minimizing L2-regularized binary cross-entropy loss using the L-BFGS solver:

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^N [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)] + \frac{1}{2C} \|\mathbf{w}\|_2^2$$

The baseline configuration sets $C = 1.0$ and $\text{max_iter} = 1000$.

3) *Random Forest Classifier*: Random Forest is an ensemble of decision trees trained on bootstrap samples of the training data. For a forest of $B = 100$ trees, each tree T_b splits node features using a randomly selected subset of features $\sqrt{V} \approx 70$. Node splitting minimizes Gini Impurity:

$$I_G(p) = 1 - \sum_{k=0}^1 p_k^2$$

The final class prediction is obtained via majority voting across all 100 decision trees.

4) *Multi-Layer Perceptron (MLP)*: The MLP is a feedforward artificial neural network comprising an input layer of $V = 5,000$ units, a single hidden layer of $h = 100$ neurons with Rectified Linear Unit (ReLU) activations, and a single sigmoid output node:

$$\mathbf{h} = \max(0, \mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\hat{y} = \sigma(\mathbf{w}_2^T \mathbf{h} + b_2)$$

Network parameters are trained using the Adam optimizer with a learning rate of $\alpha = 0.001$, minimizing cross-entropy loss across a maximum of 300 epochs.

V. EXPERIMENTAL RESULTS

A. Experimental Setup and Baseline Evaluation

The dataset was split into an 80% training partition (31,278 articles) and a 20% test partition (7,820 articles) using stratified random sampling with `random_state=42`. Models were fitted on the training split and evaluated on the test split across five key metrics: Accuracy, Precision, Recall, F1-Score, and Inference Latency. Table I summarizes baseline performance across all classifiers.

As detailed in Table I, the MLP Neural Network achieves the highest baseline test accuracy at 98.41%, closely followed by Random Forest (98.36%) and Logistic Regression (98.24%). KNN yields lower baseline accuracy (88.29%) due to high-dimensional distance concentration effects.

Figure 9 illustrates the comparative performance across all metrics, highlighting the strong performance of parametric models.

B. Confusion Matrix Analysis

Fig. 8 provides detailed confusion matrices for the test predictions:

- **KNN ($K = 5$)**: 2,879 True Negatives (TN), 702 False Positives (FP), 214 False Negatives (FN), 4,025 True Positives (TP).
- **Logistic Regression**: 3,488 TN, 93 FP, 45 FN, 4,194 TP.
- **Random Forest**: 3,493 TN, 88 FP, 40 FN, 4,199 TP.
- **MLP Neural Network**: 3,503 TN, 78 FP, 46 FN, 4,193 TP.

C. Training Time and Computational Latency Analysis

Model training duration on the 31,278-row TF-IDF matrix exhibits significant variation. Fig. 10 compares the wall-clock training times across models:

- **KNN**: 0.0152 seconds (lazy learner; fitting only stores the index).
- **Logistic Regression**: 0.1174 seconds.
- **Random Forest**: 14.5721 seconds.
- **MLP Neural Network**: 30.9381 seconds.

While KNN requires minimal training time, its inference latency is substantial (6.55 seconds to evaluate 7,820 samples) because it compares each test sample against all 31,278 training vectors. In contrast, Logistic Regression evaluates the test set in just 0.0195 seconds (0.0025 ms per sample), making it $2,620\times$ faster during inference than KNN.

TABLE I
MODEL EVALUATION METRICS ON 20% HOLD-OUT TEST SET (7,820 ARTICLES)

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	Inf. Time (s)	Per-Sample Latency
KNN ($K = 5$, Euclidean)	88.29	85.15	94.95	89.78	6.5500	0.8376 ms
Logistic Regression ($C = 1.0$)	98.24	97.83	98.94	98.38	0.0195	0.0025 ms
Random Forest (100 trees)	98.36	97.95	99.06	98.50	0.1631	0.0208 ms
MLP Neural Network (100-h)	98.41	98.17	98.91	98.54	0.1980	0.0253 ms
<i>Ablation / Variants:</i>						
KNN + LSA (200-d)	91.06	90.01	93.94	91.93	0.0041	0.5226 ms
LogReg (Tuned $C = 10.0$)	98.62	98.18	99.29	98.73	0.0198	0.0025 ms
KNN (Tuned Cosine, $K = 31$)	88.29	85.15	94.95	89.78	6.5467	0.8372 ms

D. Dateline Leakage Ablation Study

To evaluate the impact of publisher signatures in fake news datasets, an ablation experiment was conducted comparing baseline models against models trained after applying dateline removal regex filtering. Table II presents accuracy deltas post-filtering.

TABLE II
ACCURACY DELTA PRE- VS. POST-DATETIME REGEX REMOVAL

Model	Baseline (%)	Dateline Cleared (%)	Delta (%)
KNN	87.80	88.29	+0.49
LogReg	98.76	98.24	-0.52
Random Forest	99.68	98.36	-1.32
MLP Neural Net	98.87	98.41	-0.46

Stripping Reuters datelines (e.g., WASHINGTON (Reuters) -) reduced Random Forest accuracy by 1.32% and Logistic Regression accuracy by 0.52%. This performance shift demonstrates that models partially exploit repetitive source formatting markers rather than relying solely on pure semantic context.

E. Feature Importance Analysis

Random Forest feature importances were extracted to identify the unigrams contributing most to class separation. Fig. 11 displays the top 20 features ranked by Gini importance score.

Terms such as *said*, *trump*, *reuters*, *gop*, *hillary*, and *obama* carry the highest Gini reduction values, confirming that classification reliance centers on political figure names and formal reporting verbs.

F. Hyperparameter Optimization

Grid Search Cross-Validation (3-fold) was executed for KNN and Logistic Regression:

- **KNN Tuning:** Searching over $K \in \{3, 5, 7, 9, 11, 15, 21, 25, 31\}$, distance weights $\in \{\text{uniform}, \text{distance}\}$, and metrics $\in \{\text{euclidean}, \text{cosine}, \text{manhattan}\}$ selected $K = 31$, $\text{weights} = \text{'distance'}$, and $\text{metric} = \text{'cosine'}$. Cosine

distance significantly outperformed Manhattan distance (44.20% subsample accuracy). Fig. 12 illustrates KNN accuracy across K values.

- **Logistic Regression Tuning:** Searching over $C \in \{0.1, 1.0, 10.0\}$ and penalty $\in \{L1, L2\}$ selected $C = 10.0$ and $L2$ penalty. This raised test accuracy from 98.24% to **98.62%** (+0.38% improvement).

G. Statistical Significance and Bootstrap Analysis

To determine whether performance differences between top models are statistically significant, a McNemar's chi-squared test [16] was performed on paired test predictions from Random Forest and Logistic Regression. The test yielded a Yates' corrected chi-squared statistic of 0.5870 with a p-value of $p = 0.4436$. Because $p \geq 0.05$, we fail to reject the null hypothesis of equal error rates, confirming that Logistic Regression and Random Forest perform equivalently on this benchmark.

Furthermore, 1,000 bootstrap iterations [17] were run to compute 95% empirical confidence intervals for accuracy:

- **Logistic Regression 95% CI:** [97.94%, 98.52%]
- **Random Forest 95% CI:** [98.06%, 98.64%]


The heavy overlap between confidence intervals further supports the statistical equivalence of these two architectures.

H. Stratified 5-Fold Cross-Validation


To ensure results are not split-dependent, Stratified 5-Fold Cross-Validation was conducted across the entire 39,098-sample corpus. Table III presents the mean and standard deviation of accuracy across folds.

TABLE III
STRATIFIED 5-FOLD CROSS-VALIDATION RESULTS


Model Configuration	Mean Accuracy $\pm$ Std. Dev. (%)
KNN ($K = 31$, Cosine)	87.72 $\pm$ 0.41
Logistic Regression ($C = 10.0$, L2)	98.60 $\pm$ 0.06
Random Forest (100 trees)	98.19 $\pm$ 0.22
MLP Neural Network (100-h)	98.26 $\pm$ 0.11




(a) KNN ($K = 5$) Confusion Matrix



(b) Logistic Regression Confusion Matrix



(c) Random Forest Confusion Matrix



(d) MLP Neural Network Confusion Matrix

Fig. 8. Confusion matrices for all four classifiers evaluated on the 7,820 test set articles.




Fig. 9. Performance metric comparison (Accuracy, Precision, Recall, F1) across classifiers.




Fig. 10. Wall-clock model training time comparison (seconds) on 31,278 training samples.

The tight standard deviations ($\leq 0.41\%$) confirm consistent stability across folds.

I. Learning Curves and Convergence

Learning curves were generated by training models on increasing subsets ([20%, 40%, 60%, 80%, 100%]) of the training split. As shown in Fig. 13, the gap between training

and validation accuracy narrows progressively as dataset size increases, demonstrating convergence and confirming that models are not overfitting.




Fig. 11. Top 20 most important features extracted from the trained Random Forest classifier.




Fig. 12. KNN test accuracy sensitivity across K values using Cosine distance.




Fig. 13. Learning curves illustrating training vs. validation accuracy across training sample sizes.

J. ROC and Precision-Recall Curves

Discriminative performance was evaluated using Receiver Operating Characteristic (ROC) and Precision-Recall (PR) curves.




Fig. 14. ROC curves and Area Under Curve (AUC) metrics for all four classifiers.

Fig. 14 presents the ROC curves. Logistic Regression achieves an Area Under Curve (AUC) of 0.9987, Random Forest achieves 0.9985, MLP achieves 0.9983, and KNN achieves 0.9610.

Fig. 15 shows the Precision-Recall curves. All three parametric/ensemble models achieve Average Precision (AP) scores exceeding 0.9980, demonstrating exceptional precision-recall trade-offs.




Fig. 15. Precision-Recall curves and Average Precision (AP) scores across classifiers.

K. Out-of-Domain (OOD) Generalization Test

To evaluate model robustness against temporal topic drift, a smoke test ($n = 10$) was conducted using unseen news

samples published in 2023+. Accuracy dropped significantly across static models: KNN, Logistic Regression, and Random Forest all achieved 50.0% accuracy (equivalent to random guessing), whereas the MLP Neural Network retained 70.0% accuracy. This highlights a fundamental limitation of static vocabulary TF-IDF models when deployed on future, out-of-distribution news topics.

VI. DISCUSSION, TRADE-OFFS, AND SYSTEM ARCHITECTURE

A. Parametric vs. Non-Parametric Model Trade-offs

Selecting an optimal classifier involves balancing predictive accuracy against operational requirements:

- **Memory Footprint:** Non-parametric KNN requires storing the complete $31,278 \times 5,000$ training matrix in memory (47.30 MB in Compressed Sparse Row format, or 1.19 GB if dense). Conversely, Logistic Regression stores only 5,000 weight coefficients (~40 KB), making it exceptionally lightweight.
- **Disk Storage:** Random Forest consumes ~40 MB on disk to store 100 deep decision trees, whereas Logistic Regression occupies under 50 KB.
- **Inference Latency:** Logistic Regression processes inference queries in 0.0025 ms per sample, compared to KNN's 0.8372 ms.

B. Production System Architecture

Fig. 16 presents the recommended production system architecture for deploying the fake news detection pipeline in a streaming environment.




Fig. 16. Modular software architecture for real-time fake news inference service.

Raw text incoming from web hooks or news feeds passes through the regex cleaning pipeline, is vectorized via a cached 5,000-feature TF-IDF matrix transformer, and evaluated by the tuned Logistic Regression model.

VII. ETHICAL CONSIDERATIONS AND LIMITATIONS

Deploying automated fake news classifiers in production applications involves important ethical considerations:

- 1) **Censorship and Over-blocking:** High false-positive rates risk suppressing legitimate news content, impinging on open expression.

- 2) **Algorithmic Bias and Temporal Drift:** Models trained on static political datasets (e.g., 2016–2017 U.S. elections) risk misclassifying content from non-U.S. regions or future election cycles due to vocabulary shift.
- 3) **Human-in-the-Loop Moderation:** Automated scores should serve as decision-support alerts for human content moderators rather than autonomous censorship mechanisms.

VIII. CONCLUSION AND FUTURE WORK

This study implemented an end-to-end fake news detection framework using text processing from scratch. Across four classifier architectures evaluated on 39,098 articles, Logistic Regression ($C = 10.0$, L2) demonstrated the optimal operational balance, achieving **98.62% test accuracy** with an average per-sample inference time of **0.0025 ms**.

Future work will focus on:

- Integrating pretrained Transformer embeddings (e.g., DistilBERT) to improve out-of-domain generalization while maintaining low inference latency.
- Extending classification from binary decisions to multi-class stance detection and claim-verification frameworks.
- Incorporating network propagation metadata and multi-modal image-text cross-verification features.

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to the **Indian Institute of Computer Technology (IICT)** for providing me with the opportunity to undertake this internship and gain valuable practical experience in Artificial Intelligence, Machine Learning, Natural Language Processing, and Cybersecurity. The internship offered an excellent platform to strengthen my technical skills through hands-on projects, research-oriented learning, and real-world problem solving.

I am especially grateful to **Dr. Ashok Gopalakrishnan** for his exceptional guidance, mentorship, and unwavering support throughout the internship. His vast knowledge, practical teaching approach, and dedication to student learning made a significant impact on my understanding of machine learning concepts and their real-world applications. During the 15-day training program, he patiently addressed the questions and doubts of every student, ensuring that complex topics became easy to understand. His guidance during the implementation of the Titanic Machine Learning Project laid a strong foundation for my learning and greatly enhanced my confidence in applying machine learning techniques. Furthermore, his valuable suggestions and practical insights helped me identify appropriate datasets from Kaggle, enabling me to successfully complete both internship projects with confidence and accuracy.

I would also like to extend my heartfelt appreciation to all the faculty members, mentors, and coordinators at **IICT** for their continuous encouragement, constructive feedback, and support throughout the internship. Their collective efforts created an enriching learning environment that encouraged curiosity, innovation, and independent problem-solving.

Finally, I express my sincere gratitude to **Jecrc University**, for providing me with a strong academic foundation and continuous encouragement to pursue practical learning opportunities. The knowledge, experience, and confidence gained during this internship have significantly contributed to my academic and professional development and will serve as a strong foundation for my future career in Artificial Intelligence, Machine Learning, and Cybersecurity.

APPENDIX A

PROJECT IMPLEMENTATION AND NOTEBOOKS

The complete implementation spans four Jupyter notebooks stored under notebooks/:

- week1_data_ed.ipynb: Exploratory data analysis, duplicate purging, text length distributions, and token frequency visualizer.
- week2_feature_engineering.ipynb: Preprocessing logic, manual regex tokenization, Bag-of-Words, and TF-IDF vectorizer fitting.
- week3_model_building.ipynb: Classifier initializations, model training, hyperparameter tuning via Grid-SearchCV, and model serialization.
- week4_evaluation.ipynb: Performance metric calculation, confusion matrices, ROC/PR curves, McNemar statistical testing, and cross-validation execution.

APPENDIX B

SAMPLE TEST DATA

Table IV displays sample articles drawn from the test partition, showing raw text snippets, cleaned token sequences, true ground-truth labels, and model predictions.

APPENDIX C

CODE REFERENCE: PREPROCESSING.PY

```
"""
preprocessing.py - text cleaning and manual tokenization for the fake-news
pipeline.

Rules (per project brief):
- Tokenization is manual: re.findall(r'\b[a-z]+\b', text) - no nltk.
  word_tokenize.
- Stopword list comes from NLTK (list lookup only, not the tokenizer).
- Vectorization (TF-IDF / BoW) lives in features.py, not here.
"""

import re
import nltk
import pandas as pd

# Download stopwords corpus on first run (no-op if already cached).
nltk.download("stopwords", quiet=True)
from nltk.corpus import stopwords

STOPWORDS: set[str] = set(stopwords.words("english"))

# Regex pattern for URL removal - covers http/https and bare www. links.
_URL_RE = re.compile(r"https?://\S+|www\.\S+")

# After lowercasing and URL removal, keep only lowercase letters (word tokens).
_TOKEN_RE = re.compile(r"\b[a-z]+\b")

# Regex to match Reuters datelines like 'Washington (Reuters) -' at the start
# of text
_REUTERS_DATELINE_RE = re.compile(r"^[a-z0-9\s,./#&-]{1,50}?[s]*(Reuters)\s"
*["--\s]*")

def strip_reuters_dateline(text: str) -> str:
    """
    Strips Reuters dateline (e.g. 'Washington (Reuters) -') from the start of
    the lowercased text.
    """
    if not isinstance(text, str):
        return ""
    text_lower = text.lower()
    return _REUTERS_DATELINE_RE.sub("", text_lower)
```

```
def clean_text(text: str) -> str:
    """
    Full cleaning pipeline applied to a single article string.

    Steps (order matters):
    1. Lowercase and strip Reuters datelines
    2. Strip URLs
    3. Strip punctuation and digits - keep only letter characters and spaces
    4. Tokenize via regex word-boundary match (manual, no library tokenizer)
    5. Remove stopwords using NLTK's English list
    6. Rejoin tokens into a single cleaned string for vectorization

    Returns the cleaned string. Returns "" for null/non-string input.
    """
    if not isinstance(text, str) or not text.strip():
        return ""

    # Strip Reuters dateline from the beginning
    text = strip_reuters_dateline(text)
    text = _URL_RE.sub(" ", text)

    # Extract word tokens (only [a-z] runs - removes digits and punctuation
    # implicitly).
    tokens = _TOKEN_RE.findall(text)

    # Drop stopwords and very short tokens (length <= 1 adds no signal).
    tokens = [t for t in tokens if t not in STOPWORDS and len(t) > 1]

    return " ".join(tokens)

def tokenize(text: str) -> list[str]:
    """
    Return the token list for a pre-cleaned string.
    Used in EDA where we need per-token counts, not the joined string.
    """
    return text.split()

def build_corpus(df: pd.DataFrame, text_col: str = "text") -> pd.Series:
    """
    Apply clean_text to every row in df[text_col].
    Returns a Series of cleaned strings, aligned with df's index.
    """
    return df[text_col].fillna("").apply(clean_text)
```

APPENDIX D

CODE REFERENCE: FEATURES.PY

```
"""
features.py - Bag-of-Words and TF-IDF vectorization using scikit-learn.

Vectorizer objects are fitted on the training split only, then used to
transform both train and test sets - avoids data leakage.
"""

import numpy as np
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# Baseline parameters from the project brief's own Python skeleton.
# max_features=5000 keeps the vocabulary manageable and reduces noise from
# very rare words.
DEFAULT_MAX_FEATURES = 5000

def get_tfidf_vectorizer(max_features: int = DEFAULT_MAX_FEATURES, **kwargs) ->
    TfidfVectorizer:
    """
    Return a TfidfVectorizer configured for this project.

    TF-IDF down-weights terms that appear in almost every document
    (which carry little discriminative power) and boosts terms that are
    frequent in a document but rare across the corpus.

    Formula used internally by sklearn:
    TF(t, d) = count(t, d) (raw count)
    IDF(t) = log((1 + N) / (1 + df(t))) + 1 [sklearn smooth IDF]
    TF-IDF = TF * IDF
    """
    return TfidfVectorizer(max_features=max_features, **kwargs)

def get_bow_vectorizer(max_features: int = DEFAULT_MAX_FEATURES, **kwargs) ->
    CountVectorizer:
    """
    Return a CountVectorizer (Bag-of-Words) for this project.

    BoW represents each document as a vector of raw token counts.
    It ignores word order and document length - a simple but effective
    baseline for text classification.
    """
    return CountVectorizer(max_features=max_features, **kwargs)

def fit_transform(vectorizer, X_train, X_test):
    """
    Fit vectorizer on X_train, then transform both splits.
    Returns (X_train_vec, X_test_vec, vectorizer).
    """
    X_train_vec = vectorizer.fit_transform(X_train)
    X_test_vec = vectorizer.transform(X_test)
    return X_train_vec, X_test_vec, vectorizer

def top_features(vectorizer, X, n: int = 20) -> list[str]:
    """
```

TABLE IV
SAMPLE TEST PARTITION ARTICLES, CLEANING OUTPUTS, AND MODEL PREDICTIONS

Raw Text Snippet	Cleaned Token Sequence (Manual Preprocessing)	True Label	Prediction
“Russian submarines fire cruise missiles at Islamic state in Syria MOSCOW (Reuters) - The Russian Navy on Thursday fired seven cruise missiles at Islamic state targets in the suburbs of Syria s Deir al-Zor...”	russian, navy, thursday, fired, seven, cruise, missiles, islamic, state, targets, suburbs, syria, deir, al, zor, russian, defence...	Real (1)	Real (1)
“MICHELLE OBAMA TO HILLARY: “If You Can’t Run Your Own House...You Certainly Can’t Run The White House” [VIDEO] Remember when Mooch thought having #CrookedHillary back in the White House wasn’t in the best...”	michelle, obama, hillary, run, house, certainly, run, white, house, video, remember, mooch, thought, crookedhillary, back, white, house, best, interest, nation, yeah, neither, maybe, someone...	Fake (0)	Fake (0)
“WOW! HILLARY CAUGHT ON VIDEO In 2000 Saying She Doesn’t Like Emails Because You Can’t Hide Them From Investigators Too bad for Hillary she wasn’t actually telling the truth that time WATCH...”	wow, hillary, caught, video, saying, like, emails, hide, investigators, bad, hillary, actually, telling, truth, time, watch, hillary, clinton, saying, like, emails, hide, investigators, pic, twitter, com...	Fake (0)	Fake (0)

```

Return the top n features ranked by their total frequency or weight in the
matrix X.

For a CountVectorizer, features are ranked by raw term frequency (sum of
counts).
For a TfidfVectorizer, features are ranked by summed TF-IDF weight across
documents.
"""
feature_names = np.array(vectorizer.get_feature_names_out())
# Sum along the document axis (axis=0) to get column totals
sums = np.asarray(X.sum(axis=0)).flatten()
# Sort indices in descending order
top_indices = np.argsort(sums)[::-1][1:n]
return feature_names[top_indices].tolist()

```

```

Prints a one-line status for each model.
"""
fitted = {}
for name, model in models.items():
    print(f"Training {name}...", end=" ", flush=True)
    t0 = time.perf_counter()
    model.fit(X_train, y_train)
    elapsed = time.perf_counter() - t0
    fitted[name] = (model, elapsed)
    print(f"done ({elapsed:.1f}s)")
return fitted

def save_models(fitted: dict, out_dir: str | Path = RESULTS_DIR) -> None:
    """Serialize each fitted model to results/<name>.joblib."""
    out = Path(out_dir)
    out.mkdir(parents=True, exist_ok=True)
    for name, (model, _) in fitted.items():
        path = out / f"{name}.joblib"
        joblib.dump(model, path)
        print(f"Saved: {path}")

```

APPENDIX E

CODE REFERENCE: MODELS.PY

```

"""
models.py - training wrappers for all four classifiers.

Hyperparameters match the brief's own Python skeleton exactly as the baseline.
Any deviations from those defaults will be documented in docs/
phase3_model_report.md.
"""

import time
from pathlib import Path

import joblib
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier

RESULTS_DIR = Path("results")

def get_models() -> dict:
    """
    Return a dict of model name -> unfitted estimator, using the brief's
    baseline
    hyperparameters. random_state is fixed at 42 where applicable.
    """
    return {
        "KNN": KNeighborsClassifier(n_neighbors=5),
        "LogReg": LogisticRegression(max_iter=1000, random_state=42),
        "RandomForest": RandomForestClassifier(n_estimators=100, random_state
=42),
        "NeuralNet": MLPClassifier(hidden_layer_sizes=(100,), max_iter=300,
random_state=42),
    }

def train_all(models: dict, X_train, y_train) -> dict:
    """
    Fit each model. Returns a dict of name -> (fitted_model,
training_time_seconds).

```

```

def load_model(name: str, out_dir: str | Path = RESULTS_DIR):
    """Load a previously saved model from disk."""
    path = Path(out_dir) / f"{name}.joblib"
    return joblib.load(path)

def tune_hyperparameters(X_train, y_train) -> dict:
    """
    Perform Grid Search Cross-Validation for Logistic Regression and Random
    Forest.
    Returns a dict containing best estimators and their best parameters.
    """
    from sklearn.model_selection import GridSearchCV

    # 1. Logistic Regression tuning
    logreg = LogisticRegression(solver="liblinear", max_iter=1000, random_state
=42)
    logreg_param_grid = {
        "C": [0.1, 1.0, 10.0],
        "penalty": ["l1", "l2"]
    }
    print("Tuning Logistic Regression hyper-parameters using 3-fold CV...")
    t0 = time.perf_counter()
    grid_logreg = GridSearchCV(logreg, logreg_param_grid, cv=3, scoring="
accuracy", n_jobs=-1)
    grid_logreg.fit(X_train, y_train)
    logreg_time = time.perf_counter() - t0
    print(f"LogReg tuning complete in {logreg_time:.1f}s. Best params: {
grid_logreg.best_params_}")

    # 2. Random Forest tuning
    rf = RandomForestClassifier(random_state=42)
    rf_param_grid = {
        "n_estimators": [50, 100],
        "max_depth": [10, 20, None]
    }
    print("Tuning Random Forest hyper-parameters using 3-fold CV...")
    t0 = time.perf_counter()

```

```

grid_rf = GridSearchCV(rf, rf_param_grid, cv=3, scoring="accuracy", n_jobs
                        =-1)
grid_rf.fit(X_train, y_train)
rf_time = time.perf_counter() - t0
print(f"Random Forest tuning complete in {rf_time:.1f}s. Best params: {
      grid_rf.best_params_}")

return {
    "LogReg": grid_logreg.best_estimator_,
    "RandomForest": grid_rf.best_estimator_,
    "LogReg_params": grid_logreg.best_params_,
    "RandomForest_params": grid_rf.best_params_
}

```

APPENDIX F

CODE REFERENCE: EVALUATE.PY

```

"""
evaluate.py - metrics, confusion matrices, and comparison plots.

Used in week4_evaluation.ipynb and the IEEE report generation.
"""

import json
from pathlib import Path

import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import (
    ConfusionMatrixDisplay,
    accuracy_score,
    classification_report,
    confusion_matrix,
    f1_score,
    precision_score,
    recall_score,
)

FIGURES_DIR = Path("results/figures")
METRICS_DIR = Path("results/metrics")

def compute_metrics(y_true, y_pred, model_name: str) -> dict:
    """
    Compute accuracy, precision, recall, and F1 for a single model.
    Returns a dict that can be serialised to JSON.
    """
    return {
        "model": model_name,
        "accuracy": round(accuracy_score(y_true, y_pred), 4),
        "precision": round(precision_score(y_true, y_pred), 4),
        "recall": round(recall_score(y_true, y_pred), 4),
        "f1": round(f1_score(y_true, y_pred), 4),
    }

def save_metrics(metrics_list: list[dict], out_dir: str | Path = METRICS_DIR)
    -> None:
    out = Path(out_dir)
    out.mkdir(parents=True, exist_ok=True)
    path = out / "all_models_metrics.json"
    with open(path, "w") as f:
        json.dump(metrics_list, f, indent=2)
    print(f"Metrics saved to {path}")

def plot_confusion_matrix(y_true, y_pred, model_name: str, out_dir: str | Path
    = FIGURES_DIR) -> None:
    out = Path(out_dir)
    out.mkdir(parents=True, exist_ok=True)

    cm = confusion_matrix(y_true, y_pred)
    fig, ax = plt.subplots(figsize=(5, 4))
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Fake",
        "Real"])
    disp.plot(ax=ax, colorbar=False, cmap="Blues")
    ax.set_title(f"Confusion Matrix - {model_name}")
    plt.tight_layout()

    path = out / f"confusion_matrix_{model_name}.png"
    fig.savefig(path, dpi=150)
    plt.close(fig)
    print(f"Saved: {path}")

def plot_metrics_comparison(metrics_list: list[dict], out_dir: str | Path =
    FIGURES_DIR) -> None:
    """Bar chart comparing accuracy, precision, recall, and F1 across all
    models."""
    out = Path(out_dir)
    out.mkdir(parents=True, exist_ok=True)

    models = [m["model"] for m in metrics_list]
    metric_names = ["accuracy", "precision", "recall", "f1"]
    x = np.arange(len(models))
    width = 0.2

    fig, ax = plt.subplots(figsize=(10, 5))
    for i, metric in enumerate(metric_names):
        vals = [m[metric] for m in metrics_list]
        ax.bar(x + i * width, vals, width, label=metric.capitalize())

    ax.set_xticks(x + width * 1.5)
    ax.set_xticklabels(models)
    ax.set_ylim(0, 1.05)
    ax.set_ylabel("Score")

```

```

ax.set_title("Model Comparison - All Metrics")
ax.legend()
plt.tight_layout()

path = out / "model_comparison.png"
fig.savefig(path, dpi=150)
plt.close(fig)
print(f"Saved: {path}")

```

```

def McNemar_test(y_true, y_pred1, y_pred2) -> dict:
    """
    Perform McNemar's test on paired predictions from two models.
    y_pred1: prediction array of model 1
    y_pred2: prediction array of model 2
    """
    from scipy.stats import chi2
    y_true = np.array(y_true)
    y_pred1 = np.array(y_pred1)
    y_pred2 = np.array(y_pred2)

    # Contingency table cells:
    # b: Model 1 correct, Model 2 incorrect
    # c: Model 1 incorrect, Model 2 correct
    b = np.sum((y_pred1 == y_true) & (y_pred2 != y_true))
    c = np.sum((y_pred1 != y_true) & (y_pred2 == y_true))

    # McNemar's chi-squared with Yates' continuity correction
    total_discordant = b + c
    if total_discordant > 0:
        stat = ((abs(b - c) - 1.0) ** 2) / total_discordant
        p_val = chi2.sf(stat, 1)
    else:
        stat = 0.0
        p_val = 1.0

    return {
        "b": int(b),
        "c": int(c),
        "statistic": float(stat),
        "p_value": float(p_val)
    }

def bootstrap_accuracy_ci(y_true, y_pred, n_bootstraps: int = 1000, alpha:
    float = 0.05) -> tuple[float, float]:
    """
    Calculate the bootstrap confidence interval for model accuracy.
    """
    y_true = np.array(y_true)
    y_pred = np.array(y_pred)
    n_samples = len(y_true)

    rng = np.random.default_rng(42)
    boot_accuracies = []

    for _ in range(n_bootstraps):
        indices = rng.choice(n_samples, size=n_samples, replace=True)
        acc = np.mean(y_true[indices] == y_pred[indices])
        boot_accuracies.append(acc)

    boot_accuracies = np.sort(boot_accuracies)
    lower = boot_accuracies[int(n_bootstraps * (alpha / 2))]
    upper = boot_accuracies[int(n_bootstraps * (1.0 - alpha / 2))]

    return float(lower), float(upper)

```

APPENDIX G

CODE REFERENCE: UTILS.PY

```

"""
utils.py - shared helpers used across the pipeline.
"""

import sys
from pathlib import Path

def check_raw_data(raw_dir: str | Path = "data/raw") -> tuple[Path, Path]:
    """
    Verify Fake.csv and True.csv exist in raw_dir.
    Exits with a friendly message if either is missing.
    """
    raw = Path(raw_dir)
    fake_path = raw / "Fake.csv"
    true_path = raw / "True.csv"

    missing = []
    if not fake_path.exists():
        missing.append(str(fake_path))
    if not true_path.exists():
        missing.append(str(true_path))

    if missing:
        print("\n[ERROR] The following required data files are missing:")
        for p in missing:
            print(f"  - {p}")
        sys.exit(1)

    return fake_path, true_path

def ensure_dirs(*dirs: str | Path) -> None:
    """Create directories if they don't exist."""
    for d in dirs:
        Path(d).mkdir(parents=True, exist_ok=True)

```

REFERENCES

- [1] C. Bisaillon, “Fake and Real News Dataset,” Kaggle, 2020. [Online]. Available: <https://www.kaggle.com/datasets/clmentbisaillon/fake-and-real-news-dataset>
- [2] H. Ahmed, I. Traore, and S. Saad, “Detecting opinion spams and fake news using text classification,” *Journal of Security and Privacy*, vol. 1, no. 1, p. e9, 2018.
- [3] F. Pedregosa *et al.*, “Scikit-learn: Machine Learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [4] V. L. Rubin, Y. Chen, and N. J. Conroy, “Deception detection for news: Three types of fakes,” *Proceedings of the Association for Information Science and Technology*, vol. 52, no. 1, pp. 1–4, 2015.
- [5] S. Feng, R. Banerjee, and Y. Choi, “Syntactic stylometry for deception detection,” *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 171–175, 2012.
- [6] L. Penix-Tadsen, “Linguistic markers of online misinformation,” *Journal of Computational Linguistics*, vol. 44, no. 2, pp. 102–118, 2018.
- [7] G. Salton and C. Buckley, “Term-weighting approaches in automatic text retrieval,” *Information Processing & Management*, vol. 24, no. 5, pp. 513–523, 1988.
- [8] K. Shu, A. Sliva, S. Wang, J. Tang, and H. Liu, “Fake news detection on social media: A data mining perspective,” *ACM SIGKDD Explorations Newsletter*, vol. 19, no. 1, pp. 22–36, 2017.
- [9] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [10] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” *Proceedings of NAACL-HLT*, pp. 4171–4186, 2019.
- [11] Y. Liu *et al.*, “RoBERTa: A robustly optimized BERT pretraining approach,” *arXiv preprint arXiv:1907.11692*, 2019.
- [12] L. Wu and H. Liu, “Tracing fake news propagation in online social networks,” *ACM Transactions on Intelligent Systems and Technology*, vol. 10, no. 3, pp. 1–22, 2019.
- [13] K. Shu, S. Wang, and H. Liu, “Beyond news contents: The role of social context for fake news detection,” *Proceedings of the 12th ACM International Conference on Web Search and Data Mining (WSDM)*, pp. 312–320, 2019.
- [14] X. Zhou and R. Zafarani, “A survey of fake news: Fundamental theories, detection methods, and opportunities,” *ACM Computing Surveys (CSUR)*, vol. 53, no. 5, pp. 1–40, 2020.
- [15] B. D. Horne and S. Adali, “This just in: Fake news packs a lot in title, uses less words, and repeats itself,” *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM)*, vol. 11, no. 1, pp. 271–280, 2017.
- [16] Q. McNemar, “Note on the sampling error of the difference between correlated proportions or percentages,” *Psychometrika*, vol. 12, no. 2, pp. 153–157, 1947.
- [17] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*, CRC Press, 1993.
- [18] S. Deerwester, S. T. Dumais, G. W. Furnas, T. K. Landauer, and R. Harshman, “Indexing by latent semantic analysis,” *Journal of the American Society for Information Science*, vol. 41, no. 6, pp. 391–407, 1990.