

Notebook: week1_data_eda.ipynb

Week 1 — Data Loading, Cleaning & EDA

Project: AI-Powered Fake News Detection Using Text Classification

Phase: 1 of 4 (Days 1–7)

Notebook: week1_data_eda.ipynb

This notebook covers:

- Loading `Fake.csv` and `True.csv` from `data/raw/`
- Merging and labelling the dataset
- Basic data quality checks
- Manual text cleaning and tokenization
- EDA: class balance, article length distribution, most frequent words
- Saving the cleaned dataset to `data/processed/`

```
import sys
import os
from pathlib import Path
PROJECT_ROOT = Path(".").resolve()
if PROJECT_ROOT.name == "notebooks":
    PROJECT_ROOT = PROJECT_ROOT.parent

if str(PROJECT_ROOT) not in sys.path:
    sys.path.insert(0, str(PROJECT_ROOT))
```

```
os.chdir(PROJECT_ROOT)
print(f"Working directory: {PROJECT_ROOT}")
```

[Output]:
Working directory: .

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
```

```
from src.utils import check_raw_data, ensure_dirs
from src.preprocessing import clean_text, tokenize, build_corpus, strip_reuters_dateline
```

```
ensure_dirs("data/processed", "results/figures", "results/metrics", "results/notebook_exports")
```

```
sns.set_theme(style="whitegrid", palette="muted")
plt.rcParams["figure.dpi"] = 120
```

1. Load Raw Data

```
fake_path, true_path = check_raw_data("data/raw")
```

```
fake_df = pd.read_csv(fake_path)
true_df = pd.read_csv(true_path)
```

```
print(f"Fake articles: {len(fake_df):,}")
print(f"Real articles: {len(true_df):,}")
```

```
print(f"Total:          {len(fake_df) + len(true_df):,}")
```

```
[Output]:
```

```
Fake articles: 23,481
```

```
Real articles: 21,417
```

```
Total:          44,898
```

```
print("Fake.csv columns:", fake_df.columns.tolist())
```

```
print("True.csv columns:", true_df.columns.tolist())
```

```
print()
```

```
print("Fake.csv – first 2 rows:")
```

```
display(fake_df.head(2))
```

```
print("True.csv – first 2 rows:")
```

```
display(true_df.head(2))
```

```
[Output]:
```

```
Fake.csv columns: ['title', 'text', 'subject', 'date']
```

```
True.csv columns: ['title', 'text', 'subject', 'date']
```

```
Fake.csv – first 2 rows:
```

```
[Output]:
```

```
True.csv – first 2 rows:
```

2. Merge and Label

```
fake_df["label"] = 0 # 0 = fake
```

```
true_df["label"] = 1 # 1 = real
```

```
df = pd.concat([fake_df, true_df], ignore_index=True)
```

```
# Combine title and body into one text field for classification.
```

```
# Some articles have an empty/null text field – fallback to title only.
```

```
df["text"] = df["text"].fillna("").apply(strip_reuters_dateline)
```

```
df["text"] = df["title"].fillna("") + " " + df["text"]
```

```
df["text"] = df["text"].str.strip()
```

```
print(f"Merged dataset shape: {df.shape}")
```

```
print(f"Columns: {df.columns.tolist()}")
```

```
[Output]:
```

```
Merged dataset shape: (44898, 5)
```

```
Columns: ['title', 'text', 'subject', 'date', 'label']
```

3. Data Quality Checks

```
print("Null counts per column:")
```

```
print(df.isnull().sum())
```

```
print()
```

```
n_dupes = df.duplicated(subset="text").sum()
```

```
print(f"Duplicate rows (by text): {n_dupes}")
```

```
if n_dupes > 0:
```

```
    df = df.drop_duplicates(subset="text").reset_index(drop=True)
```

```
    print(f"After deduplication: {len(df):,} rows")
```

```
[Output]:
```

```
Null counts per column:
```

```
title      0
```

```
text       0
```

```
subject    0
```

```
date       0
```

```
label      0
```

```
dtype: int64
```

```
[Output]:
```

```
Duplicate rows (by text): 5795
```

```
[Output]:
```

```
After deduplication: 39,103 rows
```

```
# Drop rows where the combined text field is effectively empty.
```

```
empty_mask = df["text"].str.strip().eq("")
```

```
print(f"Rows with empty text: {empty_mask.sum()}")
```

```
df = df[~empty_mask].reset_index(drop=True)
```

```
print(f"Dataset after empty-text removal: {len(df):,} rows")
```

```
[Output]:
```

```
Rows with empty text: 0
```

```
Dataset after empty-text removal: 39,103 rows
```

4. EDA — Class Balance

```
class_counts = df["label"].value_counts()
```

```
print("Class distribution:")
```

```
print(f" Fake (0): {class_counts[0]:,} ({100 * class_counts[0] / len(df):.1f}%)")
```

```
print(f" Real (1): {class_counts[1]:,} ({100 * class_counts[1] / len(df):.1f}%)")
```

```
fig, ax = plt.subplots(figsize=(5, 4))
```

```
ax.bar(["Fake", "Real"], [class_counts[0], class_counts[1]], color=["#e05c5c", "#5c8ee0"])
```

```
ax.set_ylabel("Article count")
```

```
ax.set_title("Class Balance: Fake vs. Real")
```

```
for i, v in enumerate([class_counts[0], class_counts[1]]):
```

```
    ax.text(i, v + 100, f"{v:,}", ha="center", va="bottom", fontsize=9)
```

```
plt.tight_layout()
```

```
plt.savefig("results/figures/class_balance.png", dpi=150)
```

```
plt.show()
```

```
print("Saved: results/figures/class_balance.png")
```

```
[Output]:
```


```
Class distribution:
```

```
 Fake (0): 17,907 (45.8%)
```

```
 Real (1): 21,196 (54.2%)
```

```
[Output]:
```

```
Saved: results/figures/class_balance.png
```



5. EDA — Article Length Distribution

Word count before cleaning, using simple whitespace split.

```
df["word_count_raw"] = df["text"].str.split().apply(len)
```

```
print("Word count stats (raw text):")
```

```
print(df.groupby("label")["word_count_raw"].describe().round(1))
```

```
fig, ax = plt.subplots(figsize=(8, 4))
```

```
for label, color, name in [(0, "#e05c5c", "Fake"), (1, "#5c8ee0", "Real")]:
```

```
    subset = df[df["label"] == label]["word_count_raw"]
```

```
    ax.hist(subset.clip(upper=2000), bins=60, alpha=0.6, color=color, label=name)
```

```
ax.set_xlabel("Word count (clipped at 2000)")
```

```
ax.set_ylabel("Number of articles")
```

```
ax.set_title("Article Length Distribution")
```

```
ax.legend()
```

```
plt.tight_layout()
```

```
plt.savefig("results/figures/length_distribution.png", dpi=150)
```

```
plt.show()
```

```
print("Saved: results/figures/length_distribution.png")
```


[Output]:

Word count stats (raw text):

	count	mean	std	min	25%	50%	75%	max
label								
0	17907.0	428.9	357.0	2.0	279.0	385.0	514.0	8148.0
1	21196.0	394.8	273.8	4.0	157.8	368.5	533.0	5181.0

[Output]:

Saved: results/figures/length_distribution.png



6. Text Cleaning (Manual Tokenization)

Show a before/after example on one article.

```
sample_raw = df.loc[0, "text"]
```

```
sample_clean = clean_text(sample_raw)
```

```
print("=== RAW (first 400 chars) ===")
```

```
print(sample_raw[:400])
```

```
print()
```

```
print("=== CLEANED (first 400 chars) ===")
```

```
print(sample_clean[:400])
```

```
[Output]:
```

```
=== RAW (first 400 chars) ===
```

```
Donald Trump Sends Out Embarrassing New Year's Eve Message; This is Disturbing Donald Trump just couldn t wish all Americans a
```

```
=== CLEANED (first 400 chars) ===
```

```
donald trump sends embarrassing new year eve message disturbing donald trump wish americans happy new year leave instead give s
```

```
print("Cleaning full corpus – this may take a minute...")
```

```
df["clean_text"] = build_corpus(df, text_col="text")
```

Word count after cleaning (after stopwords removal and token filtering).

```
df["word_count_clean"] = df["clean_text"].str.split().apply(len)
```

```
print("Done.")
```

```
print()
```

```
print("Word count stats (cleaned text):")
```

```
print(df.groupby("label")["word_count_clean"].describe().round(1))
```

```
[Output]:
```

```
Cleaning full corpus – this may take a minute...
```

```
[Output]:
```

```
Done.
```

```
Word count stats (cleaned text):
```

	count	mean	std	min	25%	50%	75%	max
label								
0	17907.0	238.6	204.2	0.0	156.0	213.0	282.5	4961.0
1	21196.0	234.2	159.1	4.0	95.0	218.0	314.0	2360.0

Drop rows where cleaning left an empty string (very rare – mostly image-only articles).

```
empty_after_clean = df["clean_text"].str.strip().eq("")
print(f"Empty after cleaning: {empty_after_clean.sum()}")
df = df[~empty_after_clean].reset_index(drop=True)
print(f"Final dataset size: {len(df):,} rows")
```

```
[Output]:
Empty after cleaning: 5
Final dataset size: 39,098 rows
```

7. EDA — Most Frequent Words per Class

```
def top_words(df_subset: pd.Series, n: int = 20) -> list:
    all_tokens = []
    for text in df_subset:
        all_tokens.extend(text.split())
    return Counter(all_tokens).most_common(n)


fake_top = top_words(df[df["label"] == 0]["clean_text"])
real_top = top_words(df[df["label"] == 1]["clean_text"])

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

for ax, top, title, color in [
    (axes[0], fake_top, "Top 20 Words — Fake Articles", "#e05c5c"),
    (axes[1], real_top, "Top 20 Words — Real Articles", "#5c8ee0"),
]:
    words, counts = zip(*top)
    ax.barh(words[:-1], counts[:-1], color=color)
    ax.set_title(title)
    ax.set_xlabel("Frequency")

plt.tight_layout()
plt.savefig("results/figures/top_words_per_class.png", dpi=150)
plt.show()
print("Saved: results/figures/top_words_per_class.png")
```

```
[Output]:
Saved: results/figures/top_words_per_class.png
```



8. Save Processed Dataset

```
# Keep only the columns needed for downstream phases.
out_cols = ["text", "clean_text", "label", "subject"] if "subject" in df.columns else ["text", "clean_text", "label"]
df[out_cols].to_csv("data/processed/news_cleaned.csv", index=False)
print(f"Saved cleaned dataset: data/processed/news_cleaned.csv ({len(df):,} rows)")
```

```
[Output]:
```

Saved cleaned dataset: data/processed/news_cleaned.csv (39,098 rows)


9. Subject Category Distribution (if available)

```
if "subject" in df.columns:
    print("Subject distribution across full dataset:")
    print(df["subject"].value_counts())
    print()

    fig, ax = plt.subplots(figsize=(9, 4))
    df["subject"].value_counts().plot(kind="bar", ax=ax, color="#7b7bc8")
    ax.set_xlabel("Subject category")
    ax.set_ylabel("Article count")
    ax.set_title("Article Counts by Subject Category")
    ax.tick_params(axis="x", rotation=30)
    plt.tight_layout()
    plt.savefig("results/figures/subject_distribution.png", dpi=150)
    plt.show()
    print("Saved: results/figures/subject_distribution.png")
else:
    print("No 'subject' column found in this dataset version – skipping subject EDA.")
```

[Output]:
Subject distribution across full dataset:
subject
politicsNews 11216
worldnews 9980
News 9050
politics 6832
US_News 783
left-news 705
Government News 532
Name: count, dtype: int64

[Output]:
Saved: results/figures/subject_distribution.png



10. Phase 1 Summary

Print a concise summary of everything computed above — these numbers go directly into `docs/phase1\_data\_report.md`.

```
print("=" * 50)
print("PHASE 1 SUMMARY")
print("=" * 50)
```

```

print(f" Raw fake articles:          {len(fake_df):,}")
print(f" Raw real articles:         {len(true_df):,}")
print(f" Total after merge:           {len(fake_df) + len(true_df):,}")
print(f" After deduplication:           {len(df):,}")
print(f" Class balance (fake):           {100 * (df['label']==0).sum() / len(df):.1f}%")
print(f" Class balance (real):           {100 * (df['label']==1).sum() / len(df):.1f}%")
print(f" Avg words/article (raw):         {df['word_count_raw'].mean():.0f}")
print(f" Avg words/article (clean):       {df['word_count_clean'].mean():.0f}")
print("-" * 50)
print()
print("Figures saved to results/figures/")
print("Cleaned data saved to data/processed/news_cleaned.csv")

```

[Output]:

PHASE 1 SUMMARY

```

Raw fake articles:      23,481
Raw real articles:      21,417
Total after merge:      44,898
After deduplication:    39,098
Class balance (fake):    45.8%
Class balance (real):    54.2%
Avg words/article (raw):  411
Avg words/article (clean): 236

```

Figures saved to results/figures/
Cleaned data saved to data/processed/news_cleaned.csv